

Operating systems – Assignment 2

Scheduling

Lennart Jern
CS: ens16ljn

Teacher
Ahmed Aley

December 16, 2016

1 Introduction

The Linux kernel provides a number of different scheduling policies that can be used to fine tune the performance of certain applications. In this report, five different schedulers are evaluated using an artificial, CPU intensive, task. Three of the tested schedulers are “normal”, while the last two are “real-time” schedulers, meaning that they provide higher priority for their processes than the normal ones do.

The work load consists of a simple program, called `work`, that sums over a part of Grandi’s series¹ ($1 - 1 + 1 - 1 + \dots$), using a specified number of threads. The source code for `work` can be found in listing 1. Since the task is easy to parallelize, only require minimal memory access and no disk access, it should be comparable to CPU intense tasks like compression and matrix calculations.

2 Implementation (Method?)

A Bash script (`timer.sh`, listing 2) was used to time the complete task 10 times for each scheduler, for thread counts ranging from 1 to 10. See code listing 2 for the code. Additionally, each thread keeps track of the time taken from the start of its execution until it is finished, and prints this information to `stdout`, which is forwarded to data files by the bash script.

All this data was then processed by a simple Python program, `stats`, in order to calculate the median, minimum and maximum run time for each scheduler and thread count; for both the total run times and the individual thread times. This program can be found in listing 3. In addition to the statistical calculations, `stats` also produces some figures to describe the data. The figures and calculations are mostly done using the python libraries Pandas² and Matplotlib³.

It should be noted here that the processes running with real-time schedulers were run with maximum priority. Since the other schedulers does not accept a priority setting, other than the nice value, these were left untouched.

All tests were run on my personal computer with the specifications seen in table 1.

¹https://en.wikipedia.org/wiki/Grandi's_series

²<http://pandas.pydata.org/>

³<http://matplotlib.org/>

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.12-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB
GCC:	6.2.1
Bash:	4.3.43
Python:	3.5.2
Pandas:	0.18.1
Matplotlib:	1.5.3

Table 1: Test system specification.

3 Results

For the total run time of the process, there does not seem to be much of a difference between the five schedulers. This is perhaps not too surprising considering that the system was practically idle, except for the work load process, so this process naturally got all resources the system could offer.

The raw data collected can be found in appendix B.

4 Final thoughts and lessons learned




Figure 1: The median time per thread.




Figure 2: Threads max.




Figure 3: Threads min.

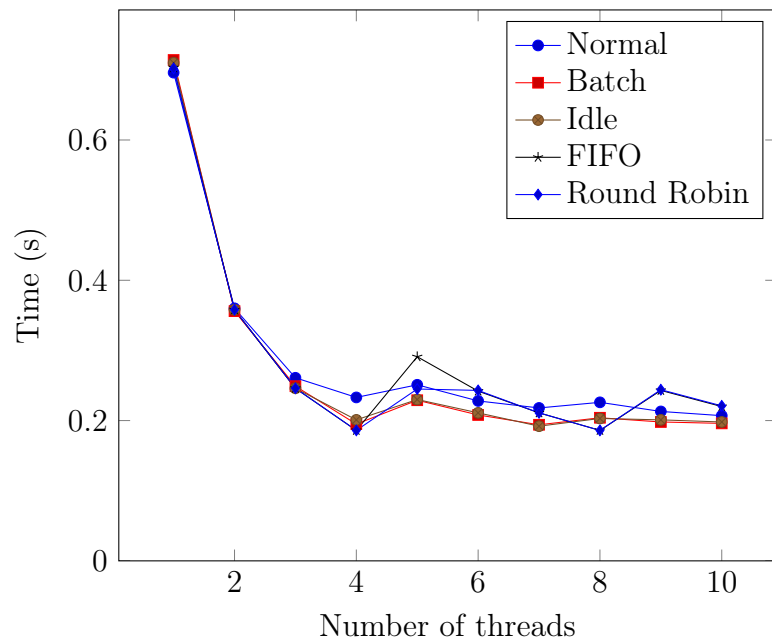


Figure 4: The median time required to finish the complete task.

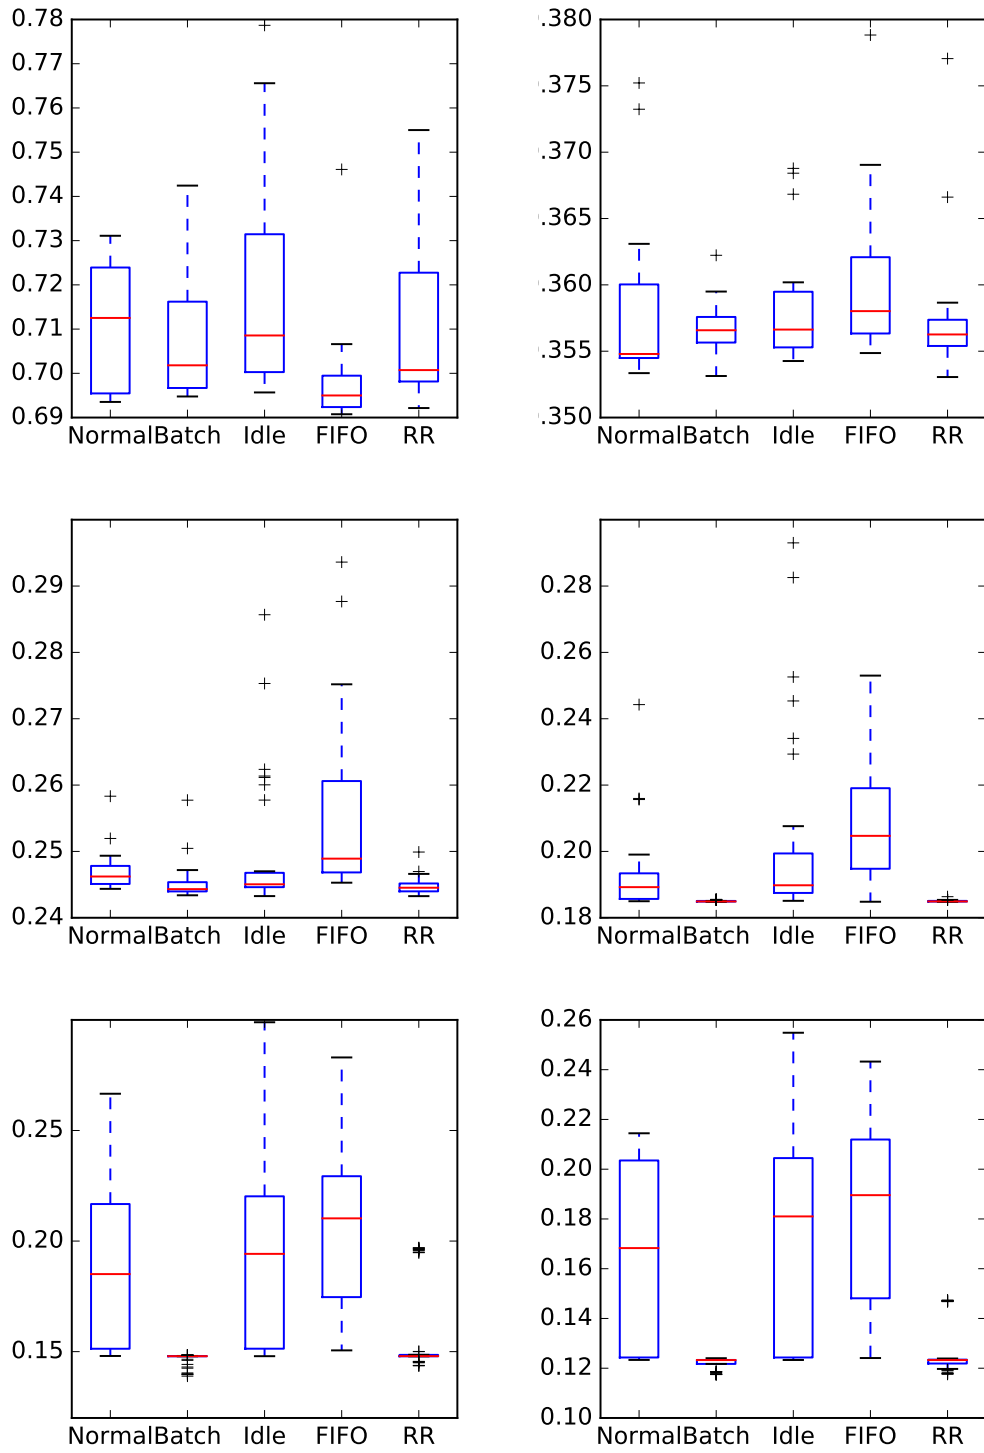


Figure 5: box1




Figure 6: box2






Figure 7: density





Figure 8: density 8

A Code listings

Listing 1: work.c

```
1  /**
2   * work.c
3   *
4   * Just a silly "do something that takes time" program.
5   * It tries to calculate the sum of Grandi's series (1-1+1-1+1-1...).
6   * As long as the length it is summing over is even the sum should always be 0.
7   *
8   * Author: Lennart Jern (ens16ljn)
9   *
10  * Call sequence: work [-p <policy>] [-j <number of jobs>]
11  * The policy is given by a single char according to this:
12  * n - Normal
13  * b - Batch
14  * i - Idle
15  * f - FIFO
16  * r - RR
17  * d - Deadline
18  */
19
20 #include <stdio.h>
21 #include <stdlib.h>
22 #include <time.h>           // timing
23 #include <errno.h>
24 #include <pthread.h>        // threading
25 #include <sys/types.h>      // pid
26 #include <unistd.h>         // pid, getopt
27 #include <linux/sched.h>    // scheduling policies
28
29 // Length of sequence to sum
30 // #define LENGTH 2147483400
31 #define LENGTH 214748340
32 // #define LENGTH 2048
33 #define ONE_OVER_BILLION 1E-9
34
35 typedef struct work_load {
36     int nworkers;
37     long data_length;
38     char scheduler;
39 } WorkLoad;
40
41 typedef struct work_packet {
42     long index;
43     long length;
44     long result;
45 } Packet;
46
47 WorkLoad *get_work_load();
48 void *work(void *data);
49 int get_grandi(int index);
50 long calculate_sum(long index, long length);
51 void run_workers(WorkLoad *wl);
52 void print_scheduler();
53 void set_scheduler(WorkLoad *wl);
54 void set_settings(WorkLoad *wl, int argc, char *argv[]);
55
56 int num_policies = 6;
57 char c_policies[] = {'n', 'b', 'i', 'f', 'r', 'd'};
```

```

58 char *str_policies[] = {"Normal", "Batch", "Idle", "FIFO", "RR", "Deadline"};
59 int policies[] = {SCHED_NORMAL, SCHED_BATCH, SCHED_IDLE, SCHED_FIFO, SCHED_RR,
    SCHED_DEADLINE};
60
61 int main(int argc, char *argv[]) {
62
63     WorkLoad *wl = get_work_load();
64
65     set_settings(wl, argc, argv);
66
67     // Set scheduler
68     set_scheduler(wl);
69
70     // Print scheduler to make sure it is set coorectly
71     print_schduler();
72
73     run_workers(wl);
74
75     free(wl);
76     printf("Done\n");
77 }
78
79 /**
80  * get_work_load - initialize the work load and return a pointer to it
81  * @return pointer to allocated memory
82  */
83 WorkLoad *get_work_load() {
84     WorkLoad *wl;
85     // Allocate memory for work load
86     wl = malloc(sizeof(WorkLoad));
87     // Initialize work load
88     wl->nworkers = 1;
89     // wl->data_length = 1073741824; // 2^30
90     wl->data_length = LENGTH;
91     return wl;
92 }
93
94 /**
95  * work - a silly attempt to calculate the limit of Grandi's series
96  * @param packet the part of the work load to work on
97  * @return nothing
98  */
99 void *work(void *packet) {
100     Packet *pkt = (Packet *)packet;
101     long sum = 0;
102
103     // Calculate time taken by a request
104     struct timespec requestStart, requestEnd;
105     clock_gettime(CLOCK_REALTIME, &requestStart);
106
107     sum = sum + calculate_sum(pkt->index, pkt->length);
108
109     clock_gettime(CLOCK_REALTIME, &requestEnd);
110
111     // Calculate time it took
112     double accum = ( requestEnd.tv_sec - requestStart.tv_sec )
113         + ( requestEnd.tv_nsec - requestStart.tv_nsec )
114         * ONE_OVER_BILLION;
115     printf( "%lf\n", accum );
116 }
117
118 /**

```

```

119  * get_grandi - calculate the i:th number of Grandi's series
120  * @param index    index of the number you want to know
121  * @return         1 if index is even, -1 otherwise
122  */
123  int get_grandi(int index) {
124      if (index % 2 == 0) {
125          return 1;
126      } else {
127          return -1;
128      }
129  }
130
131  /**
132  * calculate_sum - sum Grandi's series from index over a given length
133  * @param index    index to start from
134  * @param length   how many numbers to sum over
135  * @return         the sum
136  */
137  long calculate_sum(long index, long length) {
138      long sum = 0;
139
140      for (int i = index; i < index+length; i++) {
141          sum = sum + get_grandi(i);
142      }
143      return sum;
144  }
145
146  /**
147  * run_workers - start work threads and wait for them to finish
148  */
149  void run_workers(WorkLoad *wl) {
150      int num = wl->nworkers;
151      long total_sum = 0;
152
153      Packet *pkt[num];
154      for (int i = 0; i < num; i++) {
155          pkt[i] = malloc(sizeof(Packet));
156          if (!pkt[i]) {
157              perror("malloc");
158          }
159          pkt[i]->result = 0;
160      }
161
162      // create threads
163      pthread_t threads[num];
164      long len = wl->data_length;
165      long p_len = len / num;
166      int i;
167      for (i = 0; i < num-1; i++) {
168          pkt[i]->index = i * len / num;
169          pkt[i]->length = p_len;
170          if (pthread_create(&threads[i], NULL, work, (void *)pkt[i]) != 0) {
171              perror("Could not create thread");
172          }
173      }
174      pkt[i]->index = i * len / num;
175      pkt[i]->length = p_len;
176      work((void *)pkt[i]);
177
178      total_sum += pkt[i]->result;
179      free(pkt[i]);
180

```

```

181 // Join the threads
182 for (int i = 0; i < num-1; i++) {
183     pthread_join(threads[i], NULL);
184     total_sum += pkt[i]->result;
185     free(pkt[i]);
186 }
187
188 printf("Sum is %d\n", total_sum);
189 }
190
191 /**
192  * print_scheduler - print the current scheduler
193  * @param pid the pid of the process
194  */
195 void print_scheduler() {
196     pid_t pid = getpid();
197     int schedlr = sched_getscheduler(pid);
198
199     char *schedlr_name;
200     switch (schedlr) {
201         case SCHED_NORMAL:
202             schedlr_name = "Normal/Other";
203             break;
204         case SCHED_BATCH:
205             schedlr_name = "Batch";
206             break;
207         case SCHED_IDLE:
208             schedlr_name = "Idle";
209             break;
210         case SCHED_FIFO:
211             schedlr_name = "FIFO";
212             break;
213         case SCHED_RR:
214             schedlr_name = "RR";
215             break;
216         case SCHED_DEADLINE:
217             schedlr_name = "Deadline";
218             break;
219         default:
220             schedlr_name = "Unknown";
221     }
222     printf("Scheduler: %s\n", schedlr_name);
223 }
224
225 /**
226  * set_scheduler - update scheduler to reflect the given WorkLoad
227  * @param wl the work load
228  */
229 void set_scheduler(WorkLoad *wl) {
230     struct sched_param param;
231     pid_t pid = getpid();
232     int policy = SCHED_NORMAL;
233
234     for (int i = 0; i < num_policies; i++) {
235         if (wl->scheduler == c_policies[i]) {
236             policy = policies[i];
237             break;
238         }
239     }
240
241     // Set the priority
242     param.sched_priority = sched_get_priority_max(policy);

```

```

243
244     if (sched_setscheduler(pid, policy, &param) != 0) {
245         perror("Set scheduler");
246     }
247 }
248
249 /**
250  * set_settings - parse arguments and set the settings for the work load
251  * @param wl     the work load to update
252  * @param argc   argument count
253  * @param argv   array of arguments
254  */
255 void set_settings(WorkLoad *wl, int argc, char *argv[]) {
256     // Two possible options: j(obs) and p(policy)
257     char *optstr = "j:p:";
258     int opt;
259     char policy = 'n';
260     int num_threads = 1;
261     int policy_ok = 0;
262     int threads_ok = 0;
263
264     // Parse flags
265     while ((opt = getopt(argc, argv, optstr)) != -1) {
266         char *end;
267         switch (opt) {
268             case 'p':
269                 policy = *optarg;
270                 break;
271             case 'j':
272                 errno = 0;
273                 num_threads = strtol(optarg, &end, 10);
274                 if (errno != 0) {
275                     perror("strtol");
276                 }
277                 break;
278             default:
279                 printf("Option %c not supported\n", opt);
280         }
281     }
282
283     // Check the parsed options
284     for (int i = 0; i < num_policies; i++) {
285         if (policy == c_policies[i]) {
286             policy_ok = 1;
287             break;
288         }
289     }
290
291     if (num_threads <= 100 && num_threads > 0) {
292         threads_ok = 1;
293     }
294
295     // Set values if they are safe, or set defaults
296     if (policy_ok) {
297         wl->scheduler = policy;
298     } else {
299         wl->scheduler = 'n';
300     }
301
302     if (threads_ok) {
303         wl->nworkers = num_threads;
304     } else {

```

```

305     wl->nworkers = 1;
306 }
307 }

```

Listing 2: timer.sh

```

1  #!/bin/bash
2
3  # A timer script to measure the differences between schedulers/policies
4  #
5  # Author: Lennart Jern (ens16ljn@cs.umu.se)
6
7  for THREADS in $(seq 1 10)
8  do
9      DATA="Normal,Batch,Idle,FIFO,RR"
10     echo "Running with_${THREADS}threads"
11     # Time the commands 10 times
12     for i in $(seq 1 10)
13     do
14         LINE=""
15         # For the policies n(ormal) b(atch) and i(dle)
16         for POLICY in n b i f r
17         do
18             # Set policy and number of threads
19             FLAGS="-p$POLICY -j$THREADS"
20             COMMAND="./work_${FLAGS}_>>../data/threads$THREADS$POLICY.log"
21             # Run the command and store the time
22             t="$(sh -c "TIMEFORMAT='%5R'; time $COMMAND" 2>&1)"
23             # Build the line
24             if [ "$POLICY" = "n" ]; then
25                 LINE="$t"
26             else
27                 LINE="$LINE,$t"
28             fi
29         done
30         DATA=$DATA$'\n'$LINE
31         # A little progress report
32         echo "Run_${i}done."
33     done
34
35     # Write data to a file
36     echo "$DATA" > "../data/data$THREADS.csv"
37     chown lennart ../data/threads*
38
39 done

```

Listing 3: stats.py

```

1  """
2  stats.py
3
4  Process the data produced by timer.sh by calculating the
5  medians, max values and min values for each scheduler and thread count
6
7  Author: Lennart Jern (ens16ljn@cs.umu.se)
8  """
9
10 import pandas as pd
11 import re
12 import matplotlib.pyplot as plt

```

```

13
14 def total_stats():
15     # The data file names are of the form data<thread count>.csv
16     base = "../data/data"
17     thread_base = "../data/threads"
18     ext = ".csv"
19     header = ("Normal", "Batch", "Idle", "FIFO", "RR")
20     # Data frames to store the results in
21     med = pd.DataFrame(columns=header)          # Medians (total runtime)
22     mx = pd.DataFrame(columns=header)          # Max (total runtime)
23     mn = pd.DataFrame(columns=header)          # Min (total runtime)
24     thread_med = pd.DataFrame(columns=header)  # Medians (threads)
25     thread_mx = pd.DataFrame(columns=header)  # Max (threads)
26     thread_mn = pd.DataFrame(columns=header)  # Min (threads)
27
28     # For each number of threads
29     for i in range(1,11):
30         # Build the file name
31         f = base + str(i) + ext                # Total run times
32         thr_f = thread_base + str(i) + ext    # Thread times
33         # Read the time data
34         df = pd.read_csv(f)
35         thr_df = pd.read_csv(thr_f)
36
37         # Calculate some statistical properties
38         med.loc[i] = df.median()
39         mx.loc[i] = df.max()
40         mn.loc[i] = df.min()
41         thread_med.loc[i] = thr_df.median()
42         thread_mx.loc[i] = thr_df.max()
43         thread_mn.loc[i] = thr_df.min()
44
45         # Plot and save some nice figures
46         # Density curves for thread count 2, 4 and 8
47         if (i == 2):
48             ax = thr_df.plot.kde()
49             fig = ax.get_figure()
50             fig.savefig('density2.pdf')
51
52         if (i == 4 or i == 8):
53             ax = thr_df[["Normal", "Idle", "FIFO"]].plot.kde(figsize=(6,4))
54             fig = ax.get_figure()
55             fig.savefig("density"+str(i)+"_nif.pdf")
56             ax2 = thr_df[["Batch", "RR"]].plot.kde(figsize=(6,4))
57             fig2 = ax2.get_figure()
58             fig2.savefig("density"+str(i)+"_br.pdf")
59
60         # Box plots for all thread counts
61         ax = thr_df.plot.box(figsize=(3.5,3.5))
62         ax.set_ylabel("Time(s)")
63         fig = ax.get_figure()
64         fig.savefig("box"+str(i)+".pdf")
65
66     # Calculate ranges
67     rng = mx-mn
68     thr_rng = thread_mx-thread_mn
69
70     # Write everything to files
71     data_frames = [med, mx, mn, thread_med, thread_mx, thread_mn, rng, thr_rng]
72     base = "../data/"
73     names = ["medians", "max", "min", "thread_medians", "thread_max",
74

```

```

75         "thread_min", "range", "thread_range"]
76     for frm, name in zip(data_frames, names):
77         frm.to_csv(base+name + ext, index_label="Threads", float_format="%.5f")
78
79
80 def collect_thread_times(file_name):
81     """Read thread times from a file."""
82     f = open(file_name)
83     times = []
84     # Regular expression to find floats
85     time = re.compile("(\\d+\\.\\d+)")
86
87     for line in f:
88         match = time.match(line)
89
90         if (match):
91             t = float(match.group(1))
92             times.append(t)
93
94     return times
95
96
97 def thread_stats():
98     """Collect timing information about all threads and store in csv files"""
99     threads = [i for i in range(1, 11)]
100    schedulers = ["n", "b", "i", "f", "r"]
101    base = "../data/threads"
102    ext = ".log"
103    header=("Normal", "Batch", "Idle", "FIFO", "RR")
104
105    # Collect all times for one thread count in one file
106    for t in threads:
107        times = {key: [] for key in schedulers}
108        for s in schedulers:
109            f = get_file_name(t, s)
110            times[s] = collect_thread_times(f)
111        # Write to file
112        df = pd.DataFrame(times)
113        df.to_csv(get_csv_name(t), index=False, header=header)
114
115
116 def get_file_name(threads, scheduler):
117     """Get the file name for the data regarding <scheduler> and <threads>"""
118     base = "../data/threads"
119     ext = ".log"
120     return base + str(threads) + scheduler + ext
121
122 def get_csv_name(threads):
123     """For a specific number of threads: Get name of file to write data to"""
124     base = "../data/threads"
125     ext = ".csv"
126     return base + str(threads) + ext
127
128 thread_stats()
129 total_stats()

```

Listing 4: Makefile

```

1 all: work.c
2     gcc work.c -pthread -o work

```


B Raw data

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.696	0.714	0.710	0.703	0.702
2	0.360	0.356	0.358	0.358	0.358
3	0.261	0.249	0.246	0.246	0.246
4	0.233	0.195	0.201	0.186	0.186
5	0.251	0.229	0.230	0.291	0.245
6	0.228	0.208	0.211	0.242	0.243
7	0.218	0.194	0.192	0.211	0.211
8	0.226	0.204	0.203	0.186	0.186
9	0.213	0.198	0.201	0.243	0.244
10	0.207	0.196	0.198	0.220	0.221

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.747	0.732	0.780	0.744	0.756
2	0.379	0.380	0.370	0.363	0.378
3	0.309	0.259	0.286	0.258	0.251
4	0.254	0.246	0.294	0.190	0.188
5	0.284	0.272	0.313	0.295	0.251
6	0.260	0.215	0.256	0.248	0.247
7	0.234	0.207	0.202	0.212	0.211
8	0.307	0.241	0.213	0.188	0.190
9	0.251	0.227	0.210	0.246	0.245
10	0.241	0.212	0.202	0.222	0.222

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.691	0.694	0.696	0.696	0.693
2	0.356	0.354	0.356	0.356	0.356
3	0.250	0.247	0.245	0.244	0.245
4	0.206	0.190	0.190	0.186	0.186
5	0.231	0.216	0.221	0.288	0.244
6	0.208	0.206	0.193	0.242	0.242
7	0.203	0.190	0.189	0.211	0.211
8	0.202	0.192	0.190	0.186	0.186
9	0.205	0.190	0.195	0.242	0.243
10	0.203	0.189	0.193	0.219	0.221