

Operating systems – Assignment 2

Scheduling

Lennart Jern
CS: ens16ljn

Teacher
Ahmed Aley

December 15, 2016

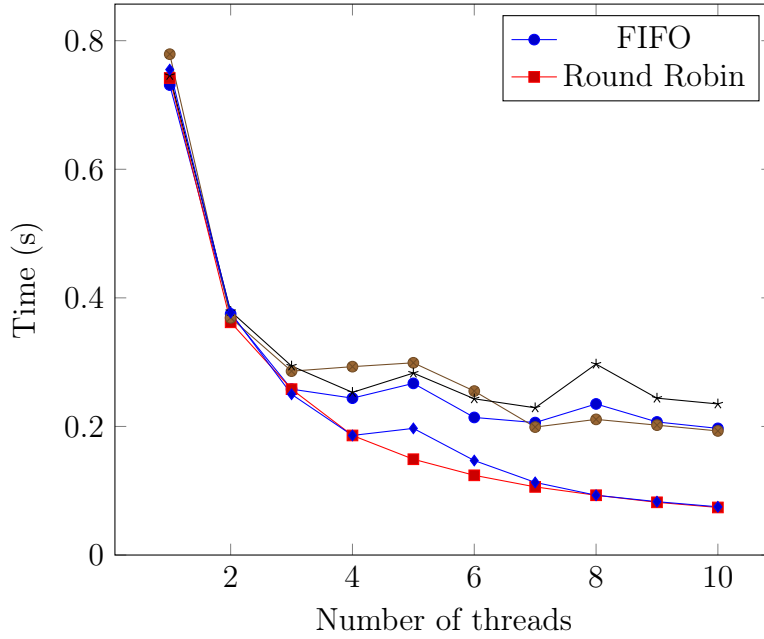


Figure 1: Threads max.

1 Introduction

The Linux kernel provides a number of different scheduling policies that can be used to fine tune the performance of certain applications. In this report, five different schedulers are evaluated using an artificial, CPU intensive, work load. Three of the tested schedulers are “normal”, while the last two are “real-time” schedulers, meaning that they provide higher priority for their processes than the normal ones do.

The work load consists of a simple program, called `work`, that sums over a part of Grandi’s series¹ ($1 - 1 + 1 - 1 + \dots$), using a specified number of threads. Since the task is easy to parallelize, only require minimal memory access and no disk access, it should be comparable to CPU intense tasks like compression and matrix calculations.

¹https://en.wikipedia.org/wiki/Grandi's_series




Figure 2: Threads min.

2 Implementation (Method?)

A Bash script (`timer.sh`) was used to collect data by timing the work load 10 times for each scheduler, for thread counts ranging from 1 to 10. See code listing 2 for the code. The data was then processed by a simple Python program in order to calculate the median, minimum and maximum run time for each scheduler and thread count. It should be noted here that the real-time schedulers were run with maximum priority. The other schedulers does not accept any priority settings.

All tests were run on my personal computer with the specifications seen in table 1.

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.12-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB

Table 1: Test system specification




Figure 3: The median time for a single thread.

3 Results

An immediate inspection of the timing data does not reveal any significant differences between the schedulers, not even between the normal and real-time ones. The median run times can be seen in figure 4 for the normal schedulers and in figure ?? for the real-time schedulers. Similarly, the maximum and minimum run times for the normal, batch and idle schedulers can be seen in figures ?? and ?? respectively, while the real-time equivalents appears in figures ?? and ??.

It is more interesting to compare the range of response times between the schedulers (fig. 5). This reveals a clear difference between the real-time and normal schedulers, where the real-time ones are clearly more predictable for two or more threads.

The raw data collected can be found in appendix B.

4 Final thoughts and lessons learned




Figure 4: The median time required to finish the complete task.

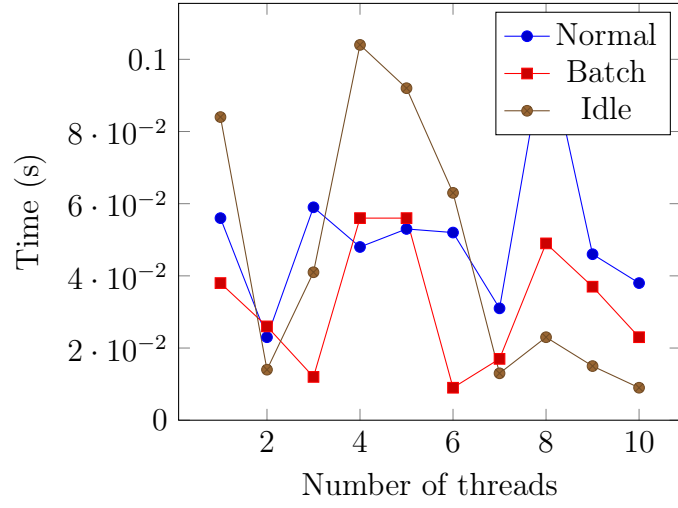
A Code listings

Listing 1: work.c

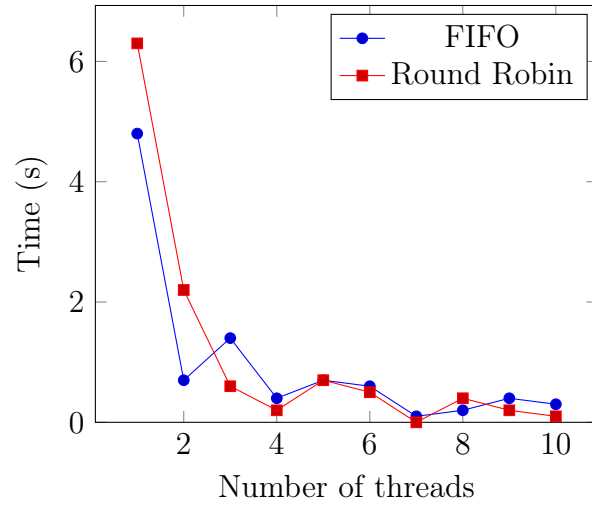
```

1  /**
2   * work.c
3   *
4   * Just a silly "do something that takes time" program.
5   * It tries to calculate the sum of Grandi's series (1-1+1-1+1-1...).
6   * As long as the length it is summing over is even the sum should always be 0.
7   *
8   * Author: Lennart Jern (ens16ljn)
9   *
10  * Call sequence: work [-p <policy>] [-j <number of jobs>]
11  * The policy is given by a single char according to this:
12  * n - Normal
13  * b - Batch
14  * i - Idle
15  * f - FIFO
16  * r - RR
17  * d - Deadline
18  */
19
20 #include <stdio.h>
21 #include <stdlib.h>
22 #include <time.h>           // timing
23 #include <errno.h>
24 #include <pthread.h>        // threading
25 #include <sys/types.h>      // pid
26 #include <unistd.h>         // pid, getopt
27 #include <linux/sched.h>    // scheduling policies

```



(a) Range for normal schedulers
 $\cdot 10^{-2}$



(b) Range for real time schedulers.

Figure 5: Range

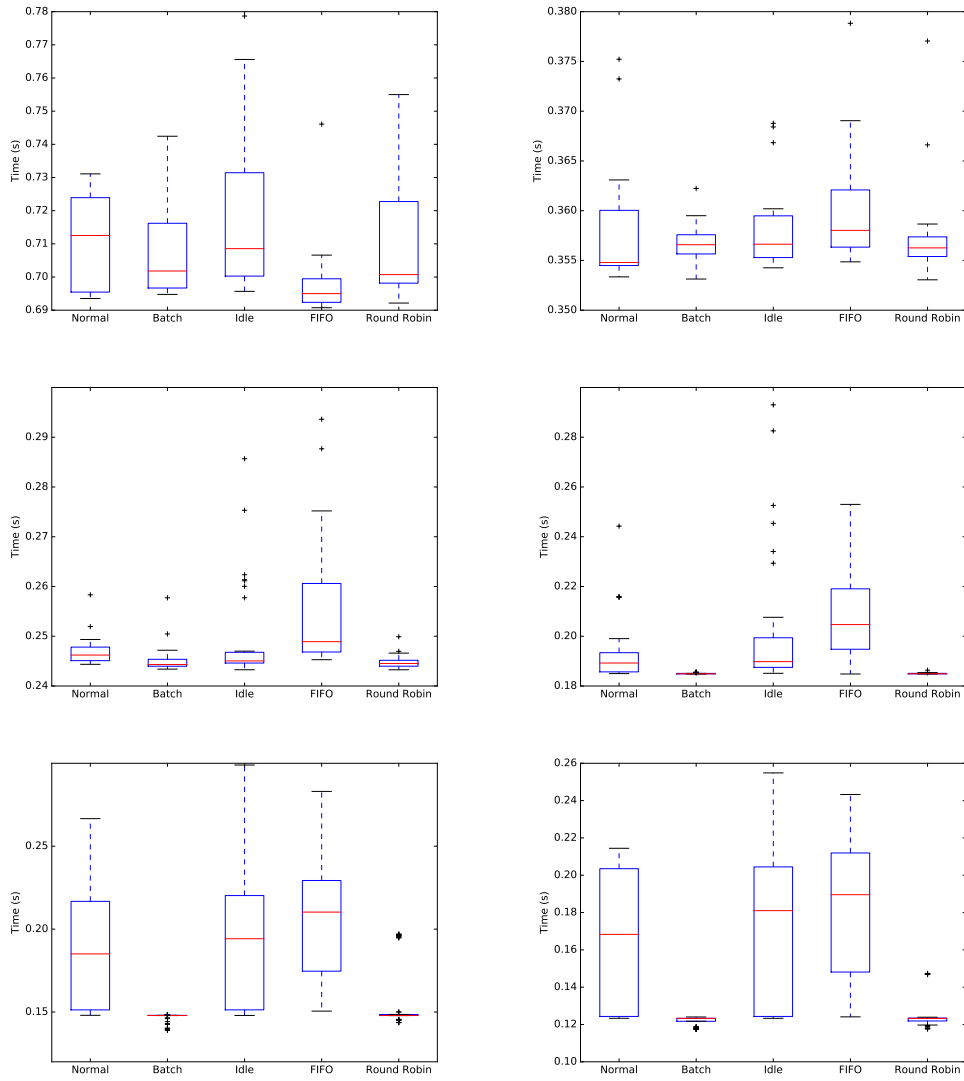


Figure 6: box1

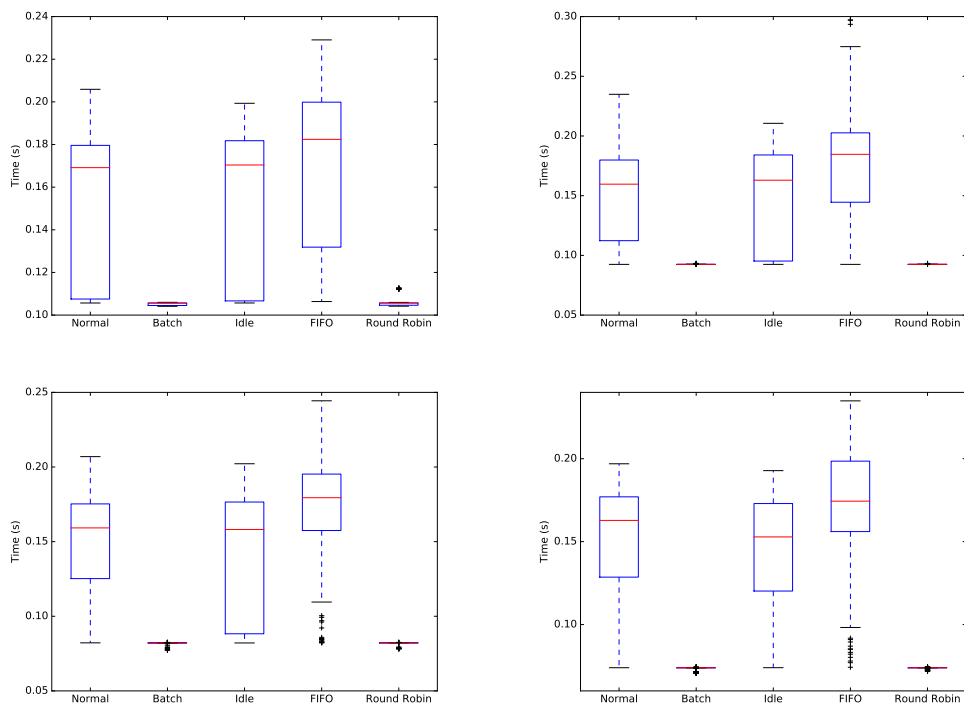


Figure 7: box2

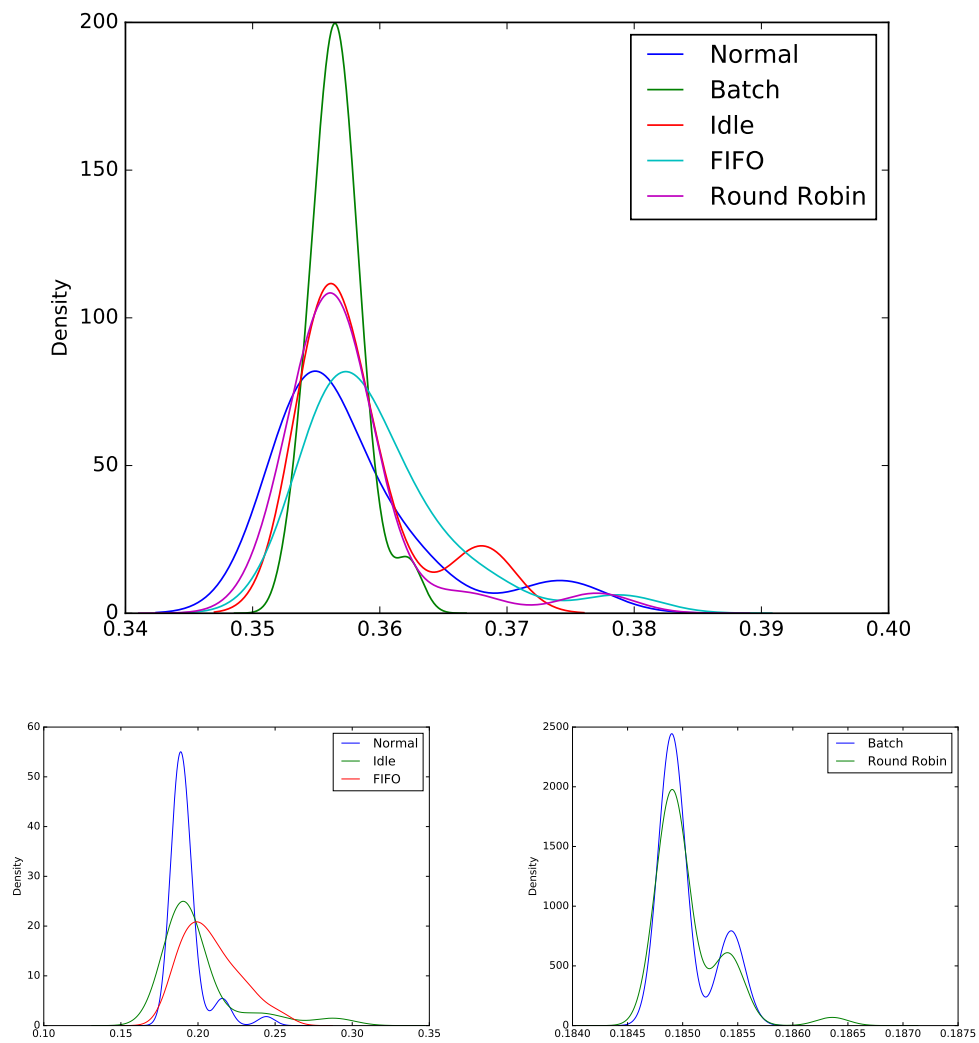


Figure 8: density




Figure 9: density 8

```

28
29 // Length of sequence to sum
30 //define LENGTH 2147483400
31 #define LENGTH 214748340
32 //define LENGTH 2048
33 #define ONE_OVER_BILLION 1E-9
34
35 typedef struct work_load {
36     int nworkers;
37     long data_length;
38     char scheduler;
39 } WorkLoad;
40
41 typedef struct work_packet {
42     long index;
43     long length;
44     long result;
45 } Packet;
46
47 WorkLoad *get_work_load();
48 void *work(void *data);
49 int get_grandi(int index);
50 long calculate_sum(long index, long length);
51 void run_workers(WorkLoad *wl);
52 void print_scheduler();
53 void set_scheduler(WorkLoad *wl);
54 void set_settings(WorkLoad *wl, int argc, char *argv[]);
55
56 int num_policies = 6;
57 char c_policies[] = {'n', 'b', 'i', 'f', 'r', 'd'};
58 char *str_policies[] = {"Normal", "Batch", "Idle", "FIFO", "RR", "Deadline"};
59 int policies[] = {SCHED_NORMAL, SCHED_BATCH, SCHED_IDLE, SCHED_FIFO, SCHED_RR,
60     SCHED_DEADLINE};
61
62 int main(int argc, char *argv[]) {
63     WorkLoad *wl = get_work_load();
64
65     set_settings(wl, argc, argv);
66
67     // Set scheduler
68     set_scheduler(wl);

```

```

69
70     // Print scheduler to make sure it is set coorectly
71     print_scheduler();
72
73     run_workers(wl);
74
75     free(wl);
76     printf("Done\n");
77 }
78
79 /**
80  * get_work_load - initialize the work load and return a pointer to it
81  * @return pointer to allocated memory
82  */
83 WorkLoad *get_work_load() {
84     WorkLoad *wl;
85     // Allocate memory for work load
86     wl = malloc(sizeof(WorkLoad));
87     // Initialize work load
88     wl->nworkers = 1;
89     // wl->data_length = 1073741824; // 2^30
90     wl->data_length = LENGTH;
91     return wl;
92 }
93
94 /**
95  * work - a silly attempt to calculate the limit of Grandi's series
96  * @param packet the part of the work load to work on
97  * @return nothing
98  */
99 void *work(void *packet) {
100     Packet *pkt = (Packet *)packet;
101     long sum = 0;
102
103     // Calculate time taken by a request
104     struct timespec requestStart, requestEnd;
105     clock_gettime(CLOCK_REALTIME, &requestStart);
106
107     sum = sum + calculate_sum(pkt->index, pkt->length);
108
109     clock_gettime(CLOCK_REALTIME, &requestEnd);
110
111     // Calculate time it took
112     double accum = ( requestEnd.tv_sec - requestStart.tv_sec )
113         + ( requestEnd.tv_nsec - requestStart.tv_nsec )
114         * ONE_OVER_BILLION;
115     printf( "%lf\n", accum );
116 }
117
118 /**
119  * get_grandi - calculate the i:th number of Grandi's series
120  * @param index index of the number you want to know
121  * @return 1 if index is even, -1 otherwise
122  */
123 int get_grandi(int index) {
124     if (index % 2 == 0) {
125         return 1;
126     } else {
127         return -1;
128     }
129 }
130

```

```

131 /**
132  * calculate_sum - sum Grandi's series from index over a given length
133  * @param index    index to start from
134  * @param length    how many numbers to sum over
135  * @return          the sum
136  */
137 long calculate_sum(long index, long length) {
138     long sum = 0;
139
140     for (int i = index; i < index+length; i++) {
141         sum = sum + get_grandi(i);
142     }
143     return sum;
144 }
145
146 /**
147  * run_workers - start work threads and wait for them to finish
148  */
149 void run_workers(WorkLoad *wl) {
150     int num = wl->nworkers;
151     long total_sum = 0;
152
153     Packet *pkt[num];
154     for (int i = 0; i < num; i++) {
155         pkt[i] = malloc(sizeof(Packet));
156         if (!pkt[i]) {
157             perror("malloc");
158         }
159         pkt[i]->result = 0;
160     }
161
162     // create threads
163     pthread_t threads[num];
164     long len = wl->data_length;
165     long p_len = len / num;
166     int i;
167     for (i = 0; i < num-1; i++) {
168         pkt[i]->index = i * len / num;
169         pkt[i]->length = p_len;
170         if (pthread_create(&threads[i], NULL, work, (void *)pkt[i]) != 0) {
171             perror("Could not create thread");
172         }
173     }
174     pkt[i]->index = i * len / num;
175     pkt[i]->length = p_len;
176     work((void *)pkt[i]);
177
178     total_sum += pkt[i]->result;
179     free(pkt[i]);
180
181     // Join the threads
182     for (int i = 0; i < num-1; i++) {
183         pthread_join(threads[i], NULL);
184         total_sum += pkt[i]->result;
185         free(pkt[i]);
186     }
187
188     printf("Sum is %d\n", total_sum);
189 }
190
191 /**
192  * print_scheduler - print the current scheduler

```

```

193  * @param pid the pid of the process
194  */
195  void print_scheduler() {
196      pid_t pid = getpid();
197      int schedlr = sched_getscheduler(pid);
198
199      char *schedlr_name;
200      switch (schedlr) {
201          case SCHED_NORMAL:
202              schedlr_name = "Normal/Other";
203              break;
204          case SCHED_BATCH:
205              schedlr_name = "Batch";
206              break;
207          case SCHED_IDLE:
208              schedlr_name = "Idle";
209              break;
210          case SCHED_FIFO:
211              schedlr_name = "FIFO";
212              break;
213          case SCHED_RR:
214              schedlr_name = "RR";
215              break;
216          case SCHED_DEADLINE:
217              schedlr_name = "Deadline";
218              break;
219          default:
220              schedlr_name = "Unknown";
221      }
222      printf("Scheduler:_%s\n", schedlr_name);
223  }
224
225  /**
226   * set_scheduler - update scheduler to reflect the given WorkLoad
227   * @param wl the work load
228   */
229  void set_scheduler(WorkLoad *wl) {
230      struct sched_param param;
231      pid_t pid = getpid();
232      int policy = SCHED_NORMAL;
233
234      for (int i = 0; i < num_policies; i++) {
235          if (wl->scheduler == c_policies[i]) {
236              policy = policies[i];
237              break;
238          }
239      }
240
241      // Set the priority
242      param.sched_priority = sched_get_priority_max(policy);
243
244      if (sched_setscheduler(pid, policy, &param) != 0) {
245          perror("Set_scheduler");
246      }
247  }
248
249  /**
250   * set_settings - parse arguments and set the settings for the work load
251   * @param wl the work load to update
252   * @param argc argument count
253   * @param argv array of arguments
254   */

```

```

255 void set_settings(WorkLoad *wl, int argc, char *argv[]) {
256     // Two possible options: j(obs) and p(policy)
257     char *optstr = "j:p:";
258     int opt;
259     char policy = 'n';
260     int num_threads = 1;
261     int policy_ok = 0;
262     int threads_ok = 0;
263
264     // Parse flags
265     while ((opt = getopt(argc, argv, optstr)) != -1) {
266         char *end;
267         switch (opt) {
268             case 'p':
269                 policy = *optarg;
270                 break;
271             case 'j':
272                 errno = 0;
273                 num_threads = strtol(optarg, &end, 10);
274                 if (errno != 0) {
275                     perror("strtol");
276                 }
277                 break;
278             default:
279                 printf("Option %c not supported\n", opt);
280         }
281     }
282
283     // Check the parsed options
284     for (int i = 0; i < num_policies; i++) {
285         if (policy == c_policies[i]) {
286             policy_ok = 1;
287             break;
288         }
289     }
290
291     if (num_threads <= 100 && num_threads > 0) {
292         threads_ok = 1;
293     }
294
295     // Set values if they are safe, or set defaults
296     if (policy_ok) {
297         wl->scheduler = policy;
298     } else {
299         wl->scheduler = 'n';
300     }
301
302     if (threads_ok) {
303         wl->nworkers = num_threads;
304     } else {
305         wl->nworkers = 1;
306     }
307 }

```

Listing 2: timer.sh

```

1 #!/bin/bash
2
3 # A timer script to measure the differences between schedulers/policies
4 #
5 # Author: Lennart Jern (ens16ljn@cs.umu.se)

```

```

6
7 for THREADS in $(seq 1 10)
8 do
9     DATA="Normal,Batch,Idle,FIFO,Round_Robin"
10    echo "Running with_${THREADS}_threads"
11    # Time the commands 10 times
12    for i in $(seq 1 10)
13    do
14        LINE=""
15        # For the policies n(ormal) b(atch) and i(dle)
16        for POLICY in n b i f r
17        do
18            # Set policy and number of threads
19            FLAGS="-p$POLICY-j$THREADS"
20            COMMAND="./work_${FLAGS}>>./data/threads$THREADS$POLICY.log"
21            # Run the command and store the time
22            t="$(sh-c"TIMEFORMAT='%5R'; time $COMMAND"2>&1)"
23            # Build the line
24            if [ "$POLICY" = "n" ]; then
25                LINE="$t"
26            else
27                LINE="$LINE,$t"
28            fi
29        done
30        DATA=$DATA$'\n'$LINE
31        # A little progress report
32        echo "Run_${i}_done."
33    done
34
35    # Write data to a file
36    echo "$DATA" > "../data/data$THREADS.csv"
37    chown lennart ../data/threads*
38
39 done

```

Listing 3: stats.py

```

1 """
2 stats.py
3
4 Process the data produced by timer.sh by calculating the
5 medians, max values and min values for each scheduler and thread count
6
7 Author: Lennart Jern (ens16ljn@cs.umu.se)
8 """
9
10 import pandas as pd
11 import re
12 import matplotlib.pyplot as plt
13
14 def total_stats():
15     # The data file names are of the form data<thread count>.csv
16     base = "../data/data"
17     thread_base = "../data/threads"
18     ext = ".csv"
19     header = ("Normal", "Batch", "Idle", "FIFO", "Round_Robin")
20     # Data frames to store the results in
21     med = pd.DataFrame(columns=header) # Medians (total runtime)
22     mx = pd.DataFrame(columns=header) # Max (total runtime)
23     mn = pd.DataFrame(columns=header) # Min (total runtime)
24     thread_med = pd.DataFrame(columns=header) # Medians (threads)

```

```

25 thread_mx = pd.DataFrame(columns=header)      # Max (threads)
26 thread_mn = pd.DataFrame(columns=header)      # Min (threads)
27
28 # For each number of threads
29 for i in range(1,11):
30     # Build the file name
31     f = base + str(i) + ext                    # Total run times
32     thr_f = thread_base + str(i) + ext        # Thread times
33     # Read the time data
34     df = pd.read_csv(f)
35     thr_df = pd.read_csv(thr_f)
36
37     # Calculate some statistical properties
38     med.loc[i] = df.median()
39     mx.loc[i] = df.max()
40     mn.loc[i] = df.min()
41     thread_med.loc[i] = thr_df.median()
42     thread_mx.loc[i] = thr_df.max()
43     thread_mn.loc[i] = thr_df.min()
44
45     # Plot and save some nice figures
46     # Density curves for thread count 2, 4 and 8
47     if (i == 2):
48         ax = thr_df.plot.kde()
49         fig = ax.get_figure()
50         fig.savefig('density2.pdf')
51
52     if (i == 4 or i == 8):
53         ax = thr_df[["Normal", "Idle", "FIFO"]].plot.kde()
54         fig = ax.get_figure()
55         fig.savefig("density"+str(i)+"_nif.pdf")
56         ax2 = thr_df[["Batch", "RoundRobin"]].plot.kde()
57         fig2 = ax2.get_figure()
58         fig2.savefig("density"+str(i)+"_br.pdf")
59
60     # Box plots for all thread counts
61     ax = thr_df.plot.box()
62     ax.set_ylabel("Time(s)")
63     fig = ax.get_figure()
64     fig.savefig("box"+str(i)+".pdf")
65
66
67     # Calculate ranges
68     rng = mx-mn
69     thr_rng = mx-mn
70
71     # Write everything to files
72     data_frames = [med, mx, mn, thread_med, thread_mx, thread_mn, rng, thr_rng]
73     base = "../data/"
74     names = ["medians", "max", "min", "thread_medians", "thread_max",
75             "thread_min", "range", "thread_range"]
76     for frm, name in zip(data_frames, names):
77         frm.to_csv(base+name + ext, index_label="Threads", float_format="%.5f")
78
79
80 def collect_thread_times(file_name):
81     """Read thread times from a file."""
82     f = open(file_name)
83     times = []
84     # Regular expression to find floats
85     time = re.compile("(\\d+\\.\\d+)")
86

```



```

87     for line in f:
88         match = time.match(line)
89
90         if (match):
91             t = float(match.group(1))
92             times.append(t)
93
94     return times
95
96
97 def thread_stats():
98     """Collect timing information about all threads and store in csv files"""
99     threads = [i for i in range(1, 11)]
100    schedulers = ["n", "b", "i", "f", "r"]
101    base = "../data/threads"
102    ext = ".log"
103    header=("Normal", "Batch", "Idle", "FIFO", "Round_Robin")
104
105    # Collect all times for one thread count in one file
106    for t in threads:
107        times = {key: [] for key in schedulers}
108        for s in schedulers:
109            f = get_file_name(t, s)
110            times[s] = collect_thread_times(f)
111        # Write to file
112        df = pd.DataFrame(times)
113        df.to_csv(get_csv_name(t), index=False, header=header)
114
115
116 def get_file_name(threads, scheduler):
117     """Get the file name for the data regarding <scheduler> and <threads>"""
118     base = "../data/threads"
119     ext = ".log"
120     return base + str(threads) + scheduler + ext
121
122 def get_csv_name(threads):
123     """For a specific number of threads: Get name of file to write data to"""
124     base = "../data/threads"
125     ext = ".csv"
126     return base + str(threads) + ext
127
128 thread_stats()
129 total_stats()

```

Listing 4: Makefile

```

1 all: work.c
2     gcc work.c -pthread -o work

```

B Raw data

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.696	0.714	0.710	0.703	0.702
2	0.360	0.356	0.358	0.358	0.358
3	0.261	0.249	0.246	0.246	0.246
4	0.233	0.195	0.201	0.186	0.186
5	0.251	0.229	0.230	0.291	0.245
6	0.228	0.208	0.211	0.242	0.243
7	0.218	0.194	0.192	0.211	0.211
8	0.226	0.204	0.203	0.186	0.186
9	0.213	0.198	0.201	0.243	0.244
10	0.207	0.196	0.198	0.220	0.221

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.747	0.732	0.780	0.744	0.756
2	0.379	0.380	0.370	0.363	0.378
3	0.309	0.259	0.286	0.258	0.251
4	0.254	0.246	0.294	0.190	0.188
5	0.284	0.272	0.313	0.295	0.251
6	0.260	0.215	0.256	0.248	0.247
7	0.234	0.207	0.202	0.212	0.211
8	0.307	0.241	0.213	0.188	0.190
9	0.251	0.227	0.210	0.246	0.245
10	0.241	0.212	0.202	0.222	0.222

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.691	0.694	0.696	0.696	0.693
2	0.356	0.354	0.356	0.356	0.356
3	0.250	0.247	0.245	0.244	0.245
4	0.206	0.190	0.190	0.186	0.186
5	0.231	0.216	0.221	0.288	0.244
6	0.208	0.206	0.193	0.242	0.242
7	0.203	0.190	0.189	0.211	0.211
8	0.202	0.192	0.190	0.186	0.186
9	0.205	0.190	0.195	0.242	0.243
10	0.203	0.189	0.193	0.219	0.221