

Operating systems – Assignment 2

Scheduling

Lennart Jern
CS: ens16ljn

Teacher
Ahmed Aley

December 12, 2016

1 Introduction

The Linux kernel provides a number of different scheduling policies that can be used to fine tune the performance of certain applications. In this report, five different schedulers are evaluated using an artificial, CPU intensive, work load. Three of the tested schedulers are “normal”, while the last two are “real-time” schedulers, meaning that they provide higher priority for their processes than the normal ones do.

The work load consists of a simple program, called `work`, that sums over a part of Grandi’s series¹ ($1 - 1 + 1 - 1 + \dots$), using a specified number of threads. Since the task is easy to parallelize, only require minimal memory access and no disk access, it should be comparable to CPU intense tasks like compression and matrix calculations.

2 Implementation (Method?)

A Bash script (`timer.sh`) was used to collect data by timing the work load 10 times for each scheduler, for thread counts ranging from 1 to 10. See code listing 3 for the code. The data was then processed by a simple Python program in order to calculate the median, minimum and maximum run time for each scheduler and thread count. It should be noted here that the real-time schedulers were run with maximum priority. The other schedulers does not accept any priority settings.

All tests were run on my personal computer with the specifications seen in table 1.

Component	Specification
OS:	Fedora 25
Kernel:	Linux 4.8.12-300.fc25.x86_64
CPU:	Intel Core i5-2500K CPU @ 3.7GHz
RAM:	7965MiB

Table 1: Test system specification

¹https://en.wikipedia.org/wiki/Grandi's_series


3 Results

An immediate inspection of the timing data does not reveal any significant differences between the schedulers, not even between the normal and real-time ones. The median run times can be seen in figure 1a for the normal schedulers and in figure 1b for the real-time schedulers. Similarly, the maximum and minimum run times for the normal, batch and idle schedulers can be seen in figures 2a and 3a respectively, while the real-time equivalents appears in figures 2b and 3b.


It is more interesting to compare the range of response times between the schedulers (fig. 4). This reveals a clear difference between the real-time and normal schedulers, where the real-time ones are clearly more predictable for two or more threads.

The raw data collected can be found in appendix B.

4 Final thoughts and lessons learned




(a) Median time




(b) Median time for real time schedulers.

Figure 1: The median time required to finish the task.




(a) Maximum time

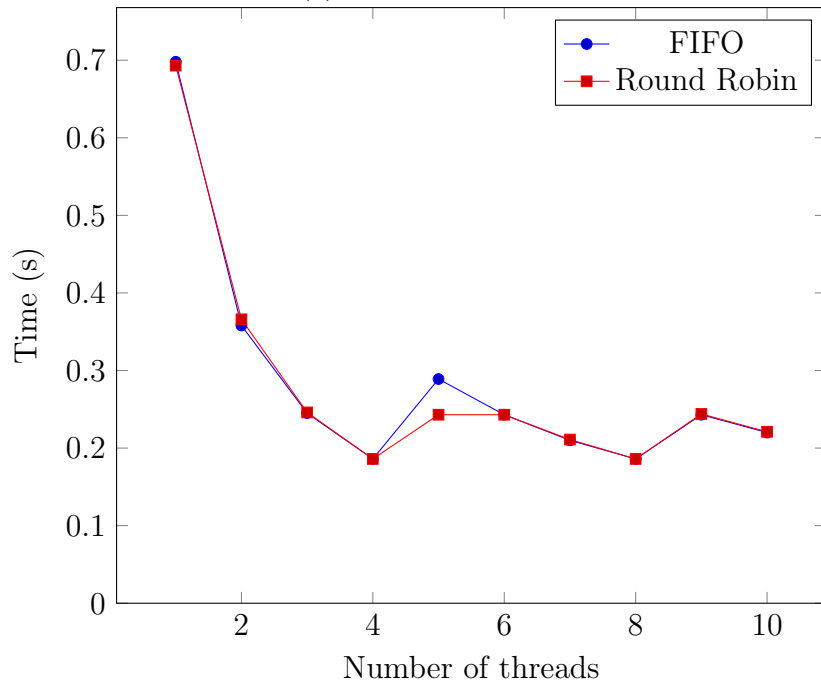


(b) Maximum time for real time schedulers.

Figure 2: The maximum run time required to complete the task.

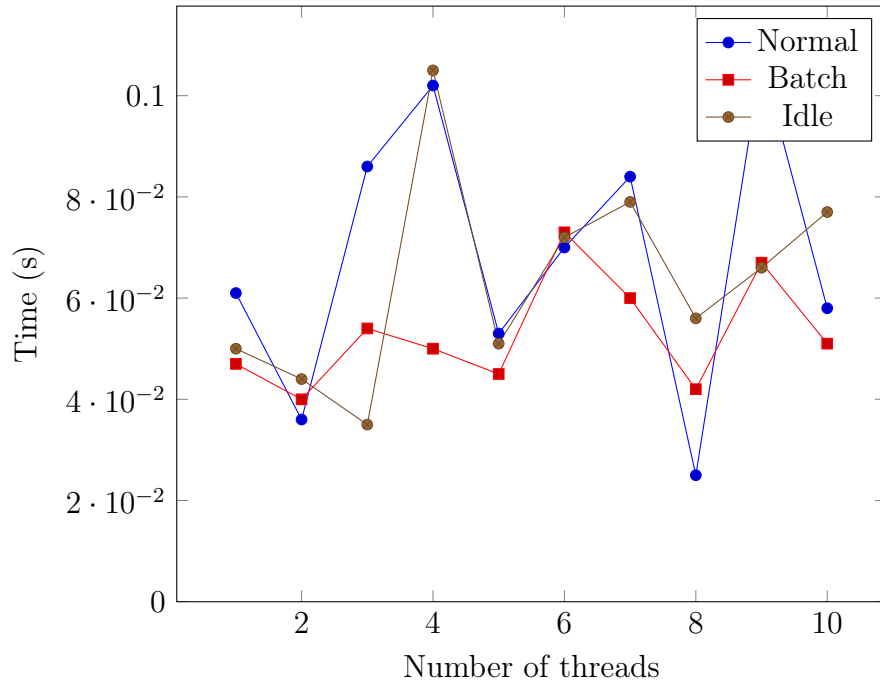


(a) Minimum time

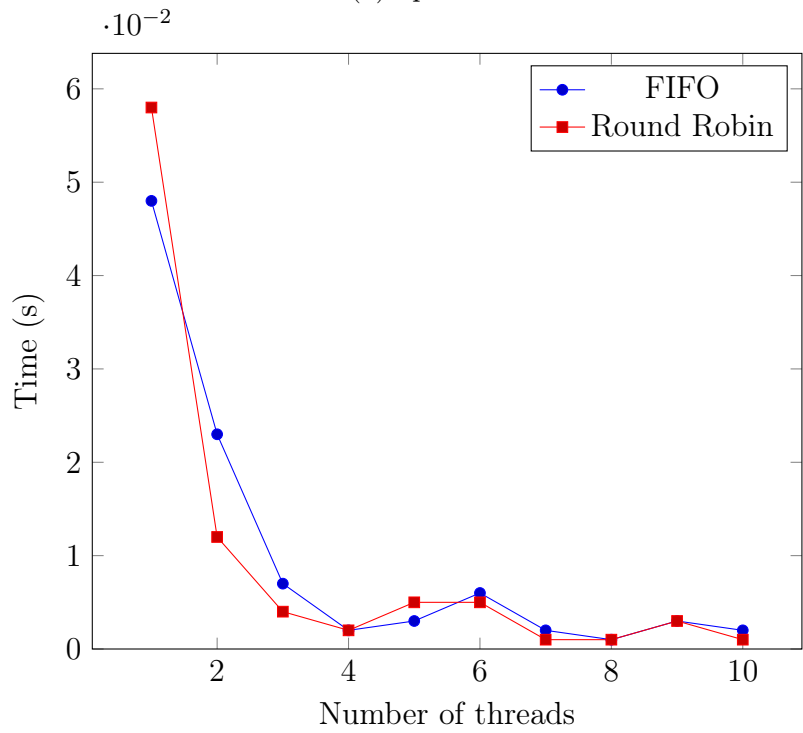


(b) Minimum time for real time schedulers.

Figure 3: The minimum run time required to complete the task.



(a) Spread



(b) Spread for real time schedulers.

Figure 4: Range

A Code listings

Listing 1: work.c

```
1  /**
2   * work.c
3   *
4   * Just a silly "do something that takes time" program.
5   * It tries to calculate the sum of Grandi's series (1-1+1-1+1-1...).
6   * As long as the length it is summing over is even the sum should always be 0.
7   *
8   * Author: Lennart Jern (ens16ljn)
9   *
10  * Call sequence: work [-p <policy>] [-j <number of jobs>]
11  * The policy is given by a single char according to this:
12  * n - Normal
13  * b - Batch
14  * i - Idle
15  * f - FIFO
16  * r - RR
17  * d - Deadline
18  */
19
20 #include <stdio.h>
21 #include <stdlib.h>
22 #include <errno.h>
23 #include <pthread.h>          // threading
24 #include <sys/types.h>        // pid
25 #include <unistd.h>           // pid, getopt
26 #include <linux/sched.h>      // scheduling policies
27
28 // Length of sequence to sum
29 // #define LENGTH 2147483400
30 #define LENGTH 214748340
31
32 typedef struct work_load {
33     int nworkers;
34     long data_length;
35     char scheduler;
36 } WorkLoad;
37
38 typedef struct work_packet {
39     long index;
40     long length;
41     long result;
42 } Packet;
43
44 WorkLoad *get_work_load();
45 void *work(void *data);
46 int get_grandi(int index);
47 long calculate_sum(long index, long length);
48 void run_workers(WorkLoad *wl);
49 void print_scheduler();
50 void set_scheduler(WorkLoad *wl);
51 void set_settings(WorkLoad *wl, int argc, char *argv[]);
52
53 int num_policies = 6;
54 char c_policies[] = {'n', 'b', 'i', 'f', 'r', 'd'};
55 char *str_policies[] = {"Normal", "Batch", "Idle", "FIFO", "RR", "Deadline"};
56 int policies[] = {SCHED_NORMAL, SCHED_BATCH, SCHED_IDLE, SCHED_FIFO, SCHED_RR,
    SCHED_DEADLINE};
```



```

57
58 int main(int argc, char *argv[]) {
59
60     WorkLoad *wl = get_work_load();
61
62     set_settings(wl, argc, argv);
63
64     // Set scheduler
65     set_scheduler(wl);
66
67     // Print scheduler to make sure it is set coorectly
68     print_scheduler();
69
70     printf("Starting workers...\n");
71     run_workers(wl);
72
73     free(wl);
74     printf("Done\n");
75 }
76
77 /**
78  * get_work_load - initialize the work load and return a pointer to it
79  * @return pointer to allocated memory
80  */
81 WorkLoad *get_work_load() {
82     WorkLoad *wl;
83     // Allocate memory for work load
84     wl = malloc(sizeof(WorkLoad));
85     // Initialize work load
86     wl->nworkers = 1;
87     // wl->data_length = 1073741824; // 2^30
88     wl->data_length = LENGTH;
89     return wl;
90 }
91
92 /**
93  * work - a silly attempt to calculate the limit of Grandi's series
94  * @param packet the part of the work load to work on
95  * @return nothing
96  */
97 void *work(void *packet) {
98     Packet *pkt = (Packet *)packet;
99
100     long sum = 0;
101
102     sum = sum + calculate_sum(pkt->index, pkt->length);
103 }
104
105 /**
106  * get_grandi - calculate the i:th number of Grandi's series
107  * @param index index of the number you want to know
108  * @return 1 if index is even, -1 otherwise
109  */
110 int get_grandi(int index) {
111     if (index % 2 == 0) {
112         return 1;
113     } else {
114         return -1;
115     }
116 }
117
118 /**

```

```

119  * calculate_sum - sum Grandi's series from index over a given length
120  * @param index    index to start from
121  * @param length    how many numbers to sum over
122  * @return          the sum
123  */
124  long calculate_sum(long index, long length) {
125      long sum = 0;
126
127      for (int i = index; i < index+length; i++) {
128          sum = sum + get_grandi(i);
129      }
130      return sum;
131  }
132
133  /**
134  * run_workers - start work threads and wait for them to finish
135  */
136  void run_workers(WorkLoad *wl) {
137      int num = wl->nworkers;
138      long total_sum = 0;
139
140      Packet *pkt[num];
141      for (int i = 0; i < num; i++) {
142          pkt[i] = malloc(sizeof(Packet));
143          if (!pkt[i]) {
144              perror("malloc");
145          }
146          pkt[i]->result = 0;
147      }
148
149      // create threads
150      pthread_t threads[num];
151      long len = wl->data_length;
152      long p_len = len / num;
153      int i;
154      for (i = 0; i < num-1; i++) {
155          pkt[i]->index = i * len / num;
156          pkt[i]->length = p_len;
157          if (pthread_create(&threads[i], NULL, work, (void *)pkt[i]) != 0) {
158              perror("Could not create thread");
159          }
160      }
161      printf("Working...\n");
162      pkt[i]->index = i * len / num;
163      pkt[i]->length = p_len;
164      work((void *)pkt[i]);
165
166      total_sum += pkt[i]->result;
167      free(pkt[i]);
168
169      // Join the threads
170      for (int i = 0; i < num-1; i++) {
171          pthread_join(threads[i], NULL);
172          total_sum += pkt[i]->result;
173          free(pkt[i]);
174      }
175
176      printf("Sum is %d\n", total_sum);
177  }
178
179  /**
180  * print_scheduler - print the current scheduler

```

```

181  * @param pid the pid of the process
182  */
183  void print_scheduler() {
184      pid_t pid = getpid();
185      int schedlr = sched_getscheduler(pid);
186
187      char *schedlr_name;
188      switch (schedlr) {
189          case SCHED_NORMAL:
190              schedlr_name = "Normal/Other";
191              break;
192          case SCHED_BATCH:
193              schedlr_name = "Batch";
194              break;
195          case SCHED_IDLE:
196              schedlr_name = "Idle";
197              break;
198          case SCHED_FIFO:
199              schedlr_name = "FIFO";
200              break;
201          case SCHED_RR:
202              schedlr_name = "RR";
203              break;
204          case SCHED_DEADLINE:
205              schedlr_name = "Deadline";
206              break;
207          default:
208              schedlr_name = "Unknown";
209      }
210      printf("Scheduler:_%s\n", schedlr_name);
211  }
212
213  /**
214   * set_scheduler - update scheduler to reflect the given WorkLoad
215   * @param wl the work load
216   */
217  void set_scheduler(WorkLoad *wl) {
218      struct sched_param param;
219      pid_t pid = getpid();
220      int policy = SCHED_NORMAL;
221
222      for (int i = 0; i < num_policies; i++) {
223          if (wl->scheduler == c_policies[i]) {
224              policy = policies[i];
225              break;
226          }
227      }
228
229      // Set the priority
230      param.sched_priority = sched_get_priority_max(policy);
231
232      if (sched_setscheduler(pid, policy, &param) != 0) {
233          perror("Set_scheduler");
234      }
235  }
236
237  /**
238   * set_settings - parse arguments and set the settings for the work load
239   * @param wl the work load to update
240   * @param argc argument count
241   * @param argv array of arguments
242   */

```

```

243 void set_settings(WorkLoad *wl, int argc, char *argv[]) {
244     // Two possible options: j(obs) and p(policy)
245     char *optstr = "j:p:";
246     int opt;
247     char policy = 'n';
248     int num_threads = 1;
249     int policy_ok = 0;
250     int threads_ok = 0;
251
252     // Parse flags
253     while ((opt = getopt(argc, argv, optstr)) != -1) {
254         char *end;
255         switch (opt) {
256             case 'p':
257                 policy = *optarg;
258                 break;
259             case 'j':
260                 errno = 0;
261                 num_threads = strtol(optarg, &end, 10);
262                 if (errno != 0) {
263                     perror("strtol");
264                 }
265                 break;
266             default:
267                 printf("Option %c not supported\n", opt);
268         }
269     }
270
271     // Check the parsed options
272     for (int i = 0; i < num_policies; i++) {
273         if (policy == c_policies[i]) {
274             policy_ok = 1;
275             break;
276         }
277     }
278
279     if (num_threads <= 100 && num_threads > 0) {
280         threads_ok = 1;
281     }
282
283     // Set values if they are safe, or set defaults
284     if (policy_ok) {
285         wl->scheduler = policy;
286     } else {
287         wl->scheduler = 'n';
288     }
289
290     if (threads_ok) {
291         wl->nworkers = num_threads;
292     } else {
293         wl->nworkers = 1;
294     }
295 }

```

Listing 2: timer.sh

```

1 #!/bin/bash
2
3 # A timer script to measure the differences between schedulers/policies
4 #
5 # Author: Lennart Jern (ens16ljn@cs.umu.se)

```

```

6
7 for THREADS in $(seq 1 10)
8 do
9     DATA="Normal,Batch,Idle,FIFO,Round_Robin"
10    echo "Running with $THREADS threads"
11    # Time the commands 10 times
12    for i in $(seq 1 10)
13    do
14        LINE=""
15        # For the policies n(ormal) b(atch) and i(dle)
16        for POLICY in n b i f r
17        do
18            # Set policy and number of threads
19            FLAGS="-p$POLICY-$THREADS"
20            COMMAND="./work_$FLAGS_/dev/null"
21            # Run the command and store the time
22            t="$(sh -c "TIMEFORMAT='%5R'; time $COMMAND" 2>&1)"
23            # Build the line
24            if [ "$POLICY" = "n" ]; then
25                LINE="$t"
26            else
27                LINE="$LINE,$t"
28            fi
29        done
30        DATA=$DATA$'\n'$LINE
31        # A little progress report
32        echo "Run $i done."
33    done
34
35    # Write data to a file
36    echo "$DATA" > "data$THREADS.csv"
37
38 done

```

Listing 3: stats.py

```

1 """
2 stats.py
3
4 Process the data produced by timer.sh by calculating the
5 medians, max values and min values for each scheduler and thread count
6
7 Author: Lennart Jern (ens16ljn@cs.umu.se)
8 """
9
10 import pandas as pd
11
12 # The data file names are of the form data<thread count>.csv
13 name = "data"
14 ext = ".csv"
15 # Data frames to store the results in
16 medians = pd.DataFrame(columns=("Normal", "Batch", "Idle", "FIFO", "Round_Robin"))
17 mx = pd.DataFrame(columns=("Normal", "Batch", "Idle", "FIFO", "Round_Robin"))
18 mn = pd.DataFrame(columns=("Normal", "Batch", "Idle", "FIFO", "Round_Robin"))
19
20 # For each number of threads
21 for i in range(1,11):
22     # Build the file name
23     f = name + str(i) + ext
24     # Read the time data
25     df = pd.read_csv(f)

```

```

26
27     # Add data to results
28     medians.loc[i] = df.median()
29     mx.loc[i] = df.max()
30     mn.loc[i] = df.min()
31
32     # Write everything to files
33     medians.to_csv("medians" + ext, index_label="Threads", float_format="%.3f")
34     mx.to_csv("max" + ext, index_label="Threads", float_format="%.3f")
35     mn.to_csv("min" + ext, index_label="Threads", float_format="%.3f")
36
37     spread = mx-mn
38     spread.to_csv("spread" + ext, index_label="Threads", float_format="%.3f")

```

Listing 4: Makefile

```

1 all: work.c
2     gcc work.c -pthread -o work

```

B Raw data

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.718	0.712	0.719	0.717	0.709
2	0.377	0.364	0.370	0.369	0.369
3	0.271	0.252	0.250	0.249	0.248
4	0.271	0.204	0.205	0.186	0.186
5	0.238	0.223	0.228	0.290	0.245
6	0.249	0.215	0.222	0.246	0.246
7	0.240	0.197	0.197	0.211	0.212
8	0.221	0.205	0.215	0.186	0.186
9	0.207	0.201	0.207	0.244	0.245
10	0.211	0.204	0.203	0.222	0.222

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.761	0.746	0.751	0.746	0.751
2	0.398	0.398	0.405	0.381	0.378
3	0.340	0.299	0.281	0.252	0.250
4	0.310	0.241	0.297	0.188	0.188
5	0.280	0.265	0.273	0.292	0.248
6	0.279	0.283	0.282	0.249	0.248
7	0.290	0.251	0.269	0.212	0.212
8	0.231	0.243	0.261	0.187	0.187
9	0.307	0.261	0.266	0.246	0.247
10	0.263	0.245	0.270	0.222	0.222

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.700	0.699	0.701	0.698	0.693
2	0.362	0.358	0.361	0.358	0.366
3	0.254	0.245	0.246	0.245	0.246
4	0.208	0.191	0.192	0.186	0.186
5	0.227	0.220	0.222	0.289	0.243
6	0.209	0.210	0.210	0.243	0.243
7	0.206	0.191	0.190	0.210	0.211
8	0.206	0.201	0.205	0.186	0.186
9	0.200	0.194	0.200	0.243	0.244
10	0.205	0.194	0.193	0.220	0.221