

Operating systems – Assignment 2

Scheduling

Lennart Jern
CS: ens16ljn

Teacher
Ahmed Aley

December 12, 2016

1 Introduction

2 Implementation

3 Testing

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.718	0.712	0.719	0.717	0.709
2	0.377	0.364	0.370	0.369	0.369
3	0.271	0.252	0.250	0.249	0.248
4	0.271	0.204	0.205	0.186	0.186
5	0.238	0.223	0.228	0.290	0.245
6	0.249	0.215	0.222	0.246	0.246
7	0.240	0.197	0.197	0.211	0.212
8	0.221	0.205	0.215	0.186	0.186
9	0.207	0.201	0.207	0.244	0.245
10	0.211	0.204	0.203	0.222	0.222

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.761	0.746	0.751	0.746	0.751
2	0.398	0.398	0.405	0.381	0.378
3	0.340	0.299	0.281	0.252	0.250
4	0.310	0.241	0.297	0.188	0.188
5	0.280	0.265	0.273	0.292	0.248
6	0.279	0.283	0.282	0.249	0.248
7	0.290	0.251	0.269	0.212	0.212
8	0.231	0.243	0.261	0.187	0.187
9	0.307	0.261	0.266	0.246	0.247
10	0.263	0.245	0.270	0.222	0.222




Figure 1: Median time

Threads	Normal	Batch	Idle	FIFO	Round Robin
1	0.700	0.699	0.701	0.698	0.693
2	0.362	0.358	0.361	0.358	0.366
3	0.254	0.245	0.246	0.245	0.246
4	0.208	0.191	0.192	0.186	0.186
5	0.227	0.220	0.222	0.289	0.243
6	0.209	0.210	0.210	0.243	0.243
7	0.206	0.191	0.190	0.210	0.211
8	0.206	0.201	0.205	0.186	0.186
9	0.200	0.194	0.200	0.243	0.244
10	0.205	0.194	0.193	0.220	0.221

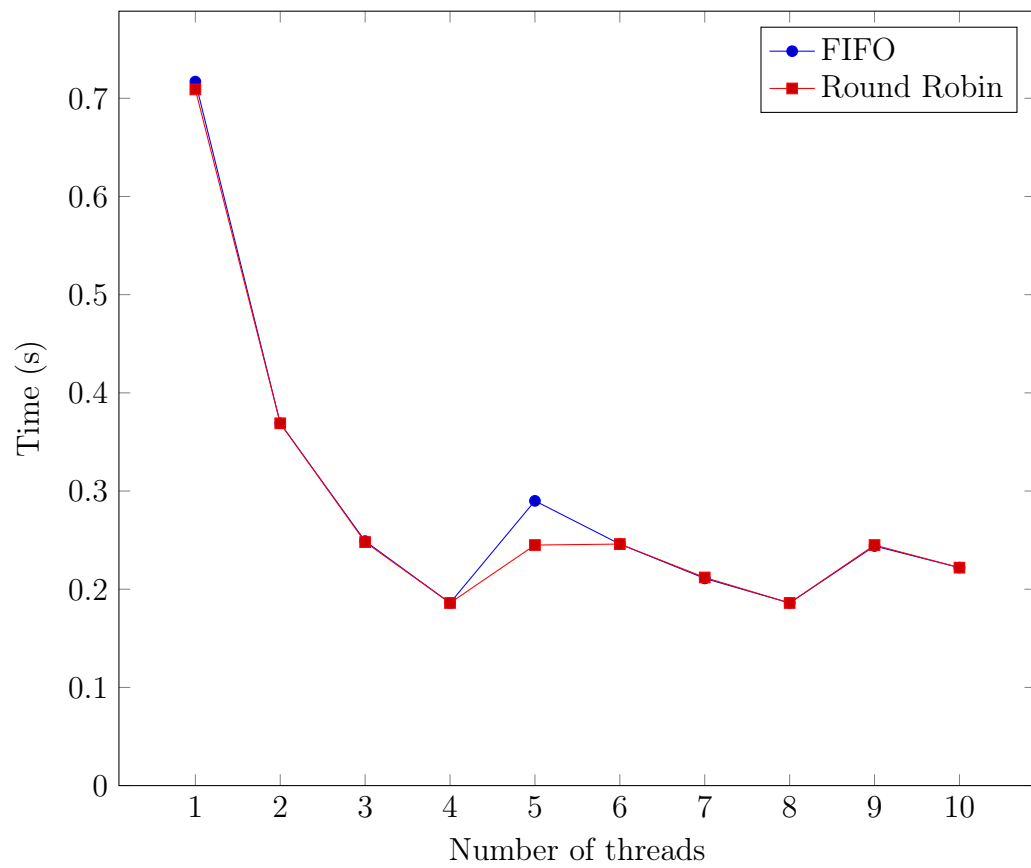


Figure 2: Median time for real time schedulers.




Figure 3: Maximum time




Figure 4: Maximum time for real time schedulers.




Figure 5: Minimum time




Figure 6: Minimum time for real time schedulers.




Figure 7: Spread

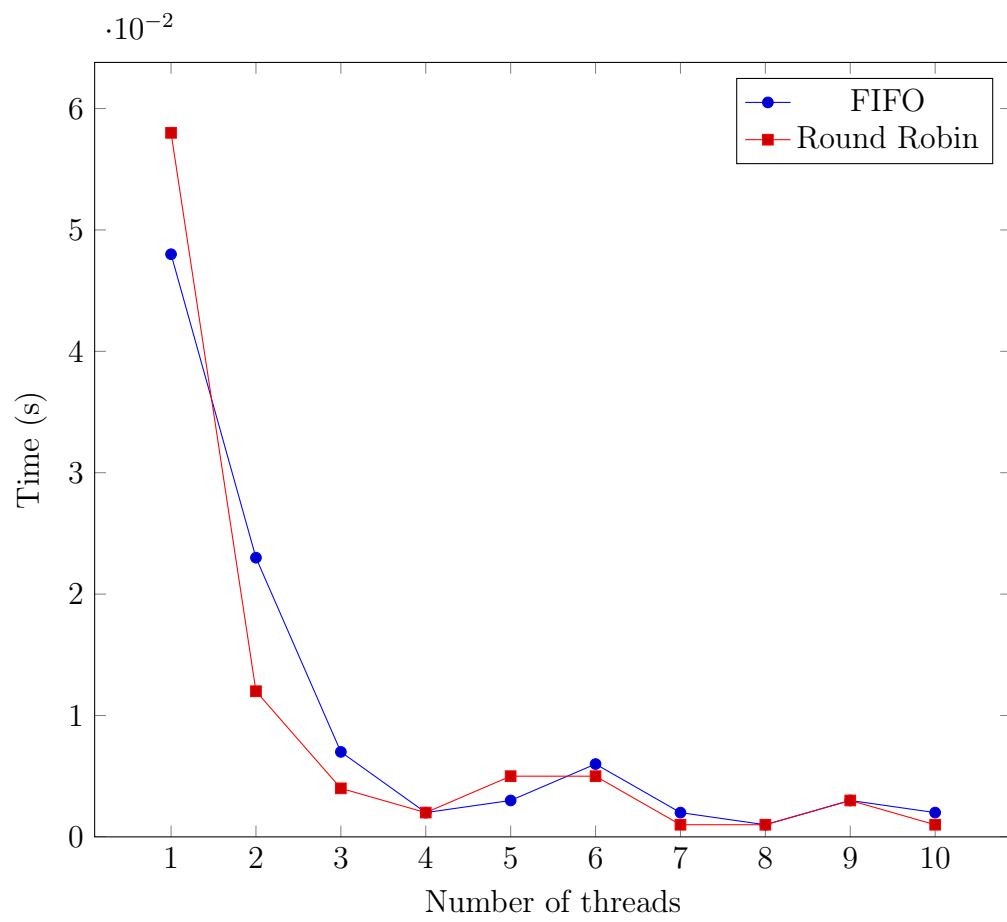


Figure 8: Spread for real time schedulers.

4 Final thoughts and lessons learned

A Code listings

```
1  /**
2   * work.c
3   *
4   * Just a silly "do something that takes time" program.
5   * It tries to calculate the sum of Grandi's series (1-1+1-1+1-1...).
6   * As long as the length it is summing over is even the sum should always be 0.
7   *
8   * Author: Lennart Jern (ens16ljn)
9   *
10  * Call sequence: work [-p <policy>] [-j <number of jobs>]
11  * The policy is given by a single char according to this:
12  * n - Normal
13  * b - Batch
14  * i - Idle
15  * f - FIFO
16  * r - RR
17  * d - Deadline
18  */
19
20 #include <stdio.h>
21 #include <stdlib.h>
22 #include <errno.h>
23 #include <pthread.h>      // threading
24 #include <sys/types.h>    // pid
25 #include <unistd.h>       // pid, getopt
26 #include <linux/sched.h>  // scheduling policies
27
28 // Length of sequence to sum
29 // #define LENGTH 2147483400
30 #define LENGTH 214748340
31
32 typedef struct work_load {
33     int nworkers;
34     long data_length;
35     char scheduler;
36 } WorkLoad;
37
38 typedef struct work_packet {
39     long index;
40     long length;
41     long result;
42 } Packet;
43
44 WorkLoad *get_work_load();
45 void *work(void *data);
46 int get_grandi(int index);
47 long calculate_sum(long index, long length);
48 void run_workers(WorkLoad *wl);
49 void print_scheduler();
50 void set_scheduler(WorkLoad *wl);
51 void set_settings(WorkLoad *wl, int argc, char *argv[]);
52
53 int num_policies = 6;
54 char c_policies[] = {'n', 'b', 'i', 'f', 'r', 'd'};
```

```

55 char *str_policies[] = {"Normal", "Batch", "Idle", "FIFO", "RR", "Deadline"};
56 int policies[] = {SCHED_NORMAL, SCHED_BATCH, SCHED_IDLE, SCHED_FIFO, SCHED_RR,
    SCHED_DEADLINE};
57
58 int main(int argc, char *argv[]) {
59     WorkLoad *wl = get_work_load();
60
61     set_settings(wl, argc, argv);
62
63     // Set scheduler
64     set_scheduler(wl);
65
66     // Print scheduler to make sure it is set coorectly
67     print_scheduler();
68
69     printf("Starting workers...\n");
70     run_workers(wl);
71
72     free(wl);
73     printf("Done\n");
74 }
75
76 /**
77  * get_work_load - initialize the work load and return a pointer to it
78  * @return pointer to allocated memory
79  */
80 WorkLoad *get_work_load() {
81     WorkLoad *wl;
82     // Allocate memory for work load
83     wl = malloc(sizeof(WorkLoad));
84     // Initialize work load
85     wl->nworkers = 1;
86     // wl->data_length = 1073741824; // 2^30
87     wl->data_length = LENGTH;
88     return wl;
89 }
90
91 /**
92  * work - a silly attempt to calculate the limit of Grandi's series
93  * @param packet the part of the work load to work on
94  * @return nothing
95  */
96 void *work(void *packet) {
97     Packet *pkt = (Packet *)packet;
98
99     long sum = 0;
100
101     sum = sum + calculate_sum(pkt->index, pkt->length);
102 }
103
104 /**
105  * get_grandi - calculate the i:th number of Grandi's series
106  * @param index index of the number you want to know
107  * @return 1 if index is even, -1 otherwise
108  */
109 int get_grandi(int index) {
110     if (index % 2 == 0) {
111         return 1;
112     } else {
113         return -1;
114     }
115 }

```

```

116 }
117
118 /**
119  * calculate_sum - sum Grandi's series from index over a given length
120  * @param index    index to start from
121  * @param length    how many numbers to sum over
122  * @return          the sum
123  */
124 long calculate_sum(long index, long length) {
125     long sum = 0;
126
127     for (int i = index; i < index+length; i++) {
128         sum = sum + get_grandi(i);
129     }
130     return sum;
131 }
132
133 /**
134  * run_workers - start work threads and wait for them to finish
135  */
136 void run_workers(WorkLoad *wl) {
137     int num = wl->nworkers;
138     long total_sum = 0;
139
140     Packet *pkt[num];
141     for (int i = 0; i < num; i++) {
142         pkt[i] = malloc(sizeof(Packet));
143         if (!pkt[i]) {
144             perror("malloc");
145         }
146         pkt[i]->result = 0;
147     }
148
149     // create threads
150     pthread_t threads[num];
151     long len = wl->data_length;
152     long p_len = len / num;
153     int i;
154     for (i = 0; i < num-1; i++) {
155         pkt[i]->index = i * len / num;
156         pkt[i]->length = p_len;
157         if (pthread_create(&threads[i], NULL, work, (void *)pkt[i]) != 0) {
158             perror("Could not create thread");
159         }
160     }
161     printf("Working...\n");
162     pkt[i]->index = i * len / num;
163     pkt[i]->length = p_len;
164     work((void *)pkt[i]);
165
166     total_sum += pkt[i]->result;
167     free(pkt[i]);
168
169     // Join the threads
170     for (int i = 0; i < num-1; i++) {
171         pthread_join(threads[i], NULL);
172         total_sum += pkt[i]->result;
173         free(pkt[i]);
174     }
175
176     printf("Sum is %d\n", total_sum);
177 }

```

```

178
179 /**
180  * print_scheduler - print the current scheduler
181  * @param pid the pid of the process
182  */
183 void print_scheduler() {
184     pid_t pid = getpid();
185     int schedlr = sched_getscheduler(pid);
186
187     char *schedlr_name;
188     switch (schedlr) {
189         case SCHED_NORMAL:
190             schedlr_name = "Normal/Other";
191             break;
192         case SCHED_BATCH:
193             schedlr_name = "Batch";
194             break;
195         case SCHED_IDLE:
196             schedlr_name = "Idle";
197             break;
198         case SCHED_FIFO:
199             schedlr_name = "FIFO";
200             break;
201         case SCHED_RR:
202             schedlr_name = "RR";
203             break;
204         case SCHED_DEADLINE:
205             schedlr_name = "Deadline";
206             break;
207         default:
208             schedlr_name = "Unknown";
209     }
210     printf("Scheduler: %s\n", schedlr_name);
211 }
212
213 /**
214  * set_scheduler - update scheduler to reflect the given WorkLoad
215  * @param wl the work load
216  */
217 void set_scheduler(WorkLoad *wl) {
218     struct sched_param param;
219     pid_t pid = getpid();
220     int policy = SCHED_NORMAL;
221
222     for (int i = 0; i < num_policies; i++) {
223         if (wl->scheduler == c_policies[i]) {
224             policy = policies[i];
225             break;
226         }
227     }
228
229     // Set the priority
230     param.sched_priority = sched_get_priority_max(policy);
231
232     if (sched_setscheduler(pid, policy, &param) != 0) {
233         perror("Set scheduler");
234     }
235 }
236
237 /**
238  * set_settings - parse arguments and set the settings for the work load
239  * @param wl the work load to update

```

```

240  * @param argc  argument count
241  * @param argv  array of arguments
242  */
243  void set_settings(WorkLoad *wl, int argc, char *argv[]) {
244      // Two possible options: j(obs) and p(olicy)
245      char *optstr = "j:p:";
246      int opt;
247      char policy = 'n';
248      int num_threads = 1;
249      int policy_ok = 0;
250      int threads_ok = 0;
251
252      // Parse flags
253      while ((opt = getopt(argc, argv, optstr)) != -1) {
254          char *end;
255          switch (opt) {
256              case 'p':
257                  policy = *optarg;
258                  break;
259              case 'j':
260                  errno = 0;
261                  num_threads = strtol(optarg, &end, 10);
262                  if (errno != 0) {
263                      perror("strtol");
264                  }
265                  break;
266              default:
267                  printf("Option_%c_not_supported\n", opt);
268          }
269      }
270
271      // Check the parsed options
272      for (int i = 0; i < num_policies; i++) {
273          if (policy == c_policies[i]) {
274              policy_ok = 1;
275              break;
276          }
277      }
278
279      if (num_threads <= 100 && num_threads > 0) {
280          threads_ok = 1;
281      }
282
283      // Set values if they are safe, or set defaults
284      if (policy_ok) {
285          wl->scheduler = policy;
286      } else {
287          wl->scheduler = 'n';
288      }
289
290      if (threads_ok) {
291          wl->nworkers = num_threads;
292      } else {
293          wl->nworkers = 1;
294      }
295  }

```

```

1  #!/bin/bash
2
3  # A timer script to measure the differences between schedulers/policies
4  #

```

```

5 # Author: Lennart Jern (ens16ljn@cs.umu.se)
6
7 for THREADS in $(seq 1 10)
8 do
9     DATA="Normal,Batch,Idle,FIFO,RoundRobin"
10    echo "Running with ${THREADS} threads"
11    # Time the commands 10 times
12    for i in $(seq 1 10)
13    do
14        LINE=""
15        # For the policies n(ormal) b(atch) and i(dle)
16        for POLICY in n b i f r
17        do
18            # Running with 4 threads/jobs
19            FLAGS="-p$POLICY -j$THREADS"
20            COMMAND="./work_${FLAGS}_/dev/null"
21            # Run the command and store the time
22            t="$(sh -c "${TIMEFORMAT='%5R'; time $COMMAND"}_2>&1)"
23            # Build the line
24            if [ "$POLICY" = "n" ]; then
25                LINE="$t"
26            else
27                LINE="$LINE,$t"
28            fi
29        done
30        DATA=$DATA$'\n'$LINE
31        # A little progress report
32        echo "Run_${i}_done."
33    done
34
35    # Write data to a file
36    echo "$DATA" > "data${THREADS}.csv"
37
38 done

```

B Raw data